

# interference

*documentation & cookbook*

current revision: 2020.1

## Contents

Concepts & features .....	2
Quick Start Application .....	3
Service as standalone .....	4
Entities .....	4
Entity class registration.....	5
Create instance of entity class .....	5
Persistent operations.....	5
Transactions isolation .....	6
Executing SQL queries.....	7
Complex event processing .....	9
Distributed persistent model.....	11
Distribute rules .....	11
@DistributedId annotation.....	12
Distributed locks .....	12
Fault Tolerance .....	12
SQL horizontal-scaling queries.....	13
Configuration parameters.....	13
Cluster configuration rules .....	15
Persistent data types .....	15

## Concepts & features

Interference is a simple distributed persistent layer for java applications:

- supports Base JPA annotations
- supports SQL queries
- supports transactions
- supports complex event processing (CEP) and based streaming SQL
- runs in the same JVM with local application
- can be used as a local or distributed SQL database
- can be used as persistent layer for a distributed application
- operates with simple objects (POJOs), annotated primarily according to the JPA specification
- allows you to make changes to data and query data from any node included in the cluster
- does not contain any coordination nodes and does not require the launch of any additional coordinators. All cluster nodes are equivalent
- note: interference is not a RDBMS in the classical sense, and it does not support ddl operations (the table structure is created on the basis of @Entity class JPA-compatible annotations) and classical RDBMS features, like foreign keys. Standard dml operations, like UPDATE, DELETE for bulk changes and check constraints are planned to be implemented in future versions.

## Quick Start Application

The *interference-test* application shows example of using the basic interference use cases. Before starting and using, read the following documentation.

Consider a basic example when the interference service used as a local persistent layer of the application and runs in the same JVM with the application.

To get started with interference, you need to include the *interference.jar* library in your project configuration. For maven *pom.xml*, this might look like this (for more details see *pom.xml* file in *interference-test* project):

```
<dependencies>
  <dependency>
    <groupId>su.interference</groupId>
    <artifactId>interference</artifactId>
    <version>2020.1</version>
  </dependency>
  ...
</dependencies>
```

Next, specify the necessary set of keys in the project (application) settings (*jmxremote* settings is optional):

```
-Dsu.interference.config=interference.properties
-Dcom.sun.management.jmxremote
-Dcom.sun.management.jmxremote.port=8888
-Dcom.sun.management.jmxremote.local.only=false
-Dcom.sun.management.jmxremote.authenticate=false
-Dcom.sun.management.jmxremote.ssl=false
-Xms256g
-Xmn512m
-Xmx4g
-XX:MaxMetaspaceSize=256m
-XX:ParallelGCThreads=8
-XX:ConcGCThreads=4
```

To run a single local interference node, you can use the standard supplied *interference.properties* configuration. Note that file *interference.properties* should be within */config* subdirectory. Next, see the configuration section.

Then, add following code into initializing section of your application:

```
Instance instance = Instance.getInstance();
Session session = Session.getSession();
instance.startupInstance(session);
```

where *Instance* is *su.inteference.core.Instance* and *Session* is *su.interference.persistent.Session*.

## Service as standalone

This option can be used when the cluster node is used solely for the purpose of further horizontal scaling of the data retrieving mechanism:

```
java -cp interference.jar
-Dsu.interference.config=interference.properties
-Dcom.sun.management.jmxremote
-Dcom.sun.management.jmxremote.port=8888
-Dcom.sun.management.jmxremote.local.only=false
-Dcom.sun.management.jmxremote.authenticate=false
-Dcom.sun.management.jmxremote.ssl=false
-Xms256g
-Xmn512m
-Xmx4g
-XX:MaxMetaspaceSize=256m
-XX:ParallelGCThreads=8
-XX:ConcGCThreads=4
su.interference.standalone.Start
```

## Entities

For use with the interference service, the entity class must be compatible with the JPA specification, i.e. satisfy the following minimum criteria:

- be annotated with `@Entity` annotation
- contain an empty public constructor
- it is mandatory to have a unique identifier field marked with `@Id` annotation
- class should not be final

**Attention!** Unlike standard JPA entities, for all persistent fields, the `getColumnNames(Session s)` and `setColumnNames(Session s)` methods must be defined, where `columnName` is the name of the persistent field of the class, and `Session` is the class `su.interference.persistent.Session`.

The current version also supports JPA annotations `@Transient`, `@Table`, `@Column`, `@Index`, `@GeneratedValue`. Validation of field values and constraints are not supported in the current version.

Attention! The field marked with the `@Id` annotation must be filled with unique data. It is highly recommended that you always use the `@GeneratedValue` annotation in conjunction with this annotation to automatically generate unique keys.

Examples of annotated entities can be found in the *interference-test* application.

## Entity class registration

For any further operations, we need a user session:

```
Session session = Session.getSession();
```

Register method registers the Dept class in the system directory, then creates and loads the corresponding proxy class for it (register Dept entity class, as example):

```
session.registerClass(Dept.class);
```

## Create instance of entity class

All further operations of inserting, changing, or retrieving data must be performed on an object obtained with the participation of the factory method `newEntity(class)`

```
Dept dept = (Dept) s.newEntity (Dept.class);
```

## Persistent operations

The newly created object is not persistent, in order to save it in the database you need to execute the `persist()` method:

```
session.persist(dept);
```

so, the session supports the following methods:

`session.newEntity(class)` - creating a proxy instance of a user class. I repeat once again that in order to use all the features of the Interference engine, all the operations described below must be performed with the proxy object constructed by this method after registering the user class, and not with the user class instance obtained through `new`. At least in this case, transactions will not be supported, and in the current version this will most likely lead to an error.

`session.persist(object)` - inserts a newly created instance into the database or saves changes to an existing one. Here you need to understand that if the identifier field is `autocomplete` (using `GeneratedValue`), then any newly created object will be saved as new in any way (insert is executed).

To change an object (update) you need to use `find()`, then make the necessary changes to the object, then execute `persist`. If identifiers are set up by the application, then using `persist` it is possible to change (update) an object in the database both in the above way (`find - change - persist`) and using a newly created object with the necessary identifier. The changed object is locked for changes from other sessions (transactions) until `commit` or `rollback` are executed (see below for transactions).

It should be noted that for objects with an `autocomplete` identifier, NO checks are made for the existence of an existing object with such an `id`. We rely entirely on the correctness of the mechanism for generating identifiers. For mass inserts, it is recommended to use just such an approach, because it is significantly faster.

`session.remove(o)` - removes an object from the database. The object is locked for changes from other sessions (transactions) until `commit` or `rollback` (see below - transactions).

`session.find(class, id)` - returns an instance from the database by identifier

**Note:** in the current version, the following types are supported for the identifier field:

```
Int, long
java.lang.Integer, java.lang.Long
java.util.concurrent.AtomicInteger
java.util.concurrent.AtomicLong
java.lang.String
```

## Transactions isolation

Interference supports transactions for read / write operations. The default isolation level is `READ COMMITTED`, which means that all changes made in any transaction will only apply to those retrievable datasets whose retrieval started after `commit` was executed in the original transaction, regardless of the retrieval duration.

All the above methods of extracting or saving data automatically create a transaction, if it has not yet been created.

To complete the transaction and apply changes to this transaction for the remaining sessions, we perform:

```
session.commit();
```

To complete the transaction and rollback the changes, we use:

```
session.rollback();
```

In addition, the following two methods can be used:

`session.lock(o)` - the current transaction receives an object lock and creates an undo-snapshot for it. Returns the current object, i.e. similar to the find method, but with getting an object lock for subsequent possible changes. The object remains locked until commit or rollback.

`session.start()` - start of the statement - all selections in the current session will return consistent data at the time the start command is executed in the current session until commit or rollback are executed. At the same time, it must be understood that the above methods for extracting data (find, get) execute the start method automatically at each start.

It is important to understand that the Entity class instance which returned by get() or find() methods is shared (Interference does not create separate instances for each session) and supports isolation. So, changing the instance field in one session will not be visible in another session until commit is called in the first. It does not require re-retrieving of the Entity class instance.

## Executing SQL queries

SQL query is run using following method:

```
s.execute(query)
```

where query is a string constant. Execute method returns `su.interference.sql.ResultSet` object which contains set of `su.interference.proxy.GenericResult` objects which is contains result row data.

Table names are indicated as fully qualified class names, for example, `su.interference.interferencetest.Dept`.

All class and field names are case sensitive.

The field names in the `ResultSet` are the specified names or aliases, if the table alias was used in the field naming, the name in the `ResultSet` will look like the table alias + field name, for example, `ddeptName` for the `d.deptName` specified in the request:

```
String sql = "select d.deptName, e.empName, e.descript from  
su.interference.interferencetest.Dept d,  
su.interference.interferencetest.Emp e where d.deptId = e.deptId";
```

```
ResultSet rs = s.execute(sql);
```

```

Object result = (GenericResult) t.poll(s);

while (result != null) {

    Object o = result.getValueByName("ddeptName");
    result = (GenericResult) t.poll(s);

}

```

All data inside the ResultSet will be consistent at the time the SQL query starts.

SQL SELECT clauses available in 2020.1 release:

**SELECT** <Group\_fn(alias1.group\_column), > alias1.column\_name1,  
alias2.column\_name2, ...

**FROM** fully\_qualified\_class\_name1 alias1, fully\_qualified\_class\_name2  
alias2, ...

<**WHERE** Condition1 **AND/OR** Condition2 ... >

<**GROUP BY** alias1.column\_name1, alias2.column\_name2, ... >

<**ORDER BY** alias1.column\_name1, alias2.column\_name2, ... >

where ConditionN is standard SQL condition, e.g.:

- alias.num\_column\_name = 12345
- alias.string\_column\_name = 'string constant'
- alias.date\_column\_name = '01.01.2019' (use config.dateformat)
- alias.column\_name IN / NOT IN [12345, 12346, ...
- alias1.column\_name1 = alias2.column\_name2
- instead equals sign ( = ) may use <, >, <>, <=, >=

Group\_fn is one of next group function: COUNT(), SUM(), MIN(), MAX(), AVG(). SELECT and FROM clauses is mandatory. In <...> described optional clauses.



## Complex event processing

Interference supports complex event processing using `SELECT STREAM` clause in SQL statement.

The basic differences between of a streaming query and the usual one are as follows:

- the `Session.execute(...)` method returns a `StreamQueue` object, which is an implementation of `ResultSet`,
- the request is executed asynchronously until `StreamQueue.stop()` method will be called or until the application terminates,
- the `StreamQueue.poll ()` method returns all records previously inserted into the table and according to the `WHERE` condition (if exist) and continues to return newly added records,
- each `StreamQueue.poll ()` method always return next record after last polled position within the session, so that, provided that the SQL request is stopped and called again within same session, data retrieve was continued from the last fixed position, in another session data will be retrieve from begin of table,
- unlike usual, a streaming request does not support transactions and always returns actually inserted rows, regardless of the use of the `commit()` method in a session inserting data (`DIRTY READS`). However, it should be keep in mind that `rollback()` physically deletes data, so use it in this context with careful.

The simplest example is a query from a table that returns either all records or filtered by some condition. Such a request can be used to broadcast events from node to node or to generate alerts of a certain type:

```
String sql = "select stream e.empName, e.descript from  
su.interference.test.entity.Event e where e.eventType = 1";
```

```
StreamQueue q = (StreamQueue) s.execute(sql);  
while (true) {  
    Object o = q.poll();  
    ...  
}
```

### Tumbling windows

This example implements the so-called streaming aggregation and assumes that the inserted records are analyzed in a strictly defined order and for each group of such records one output record will be generated using group functions such as `AVG`, `COUNT`, `MAX`, `MIN`, `SUM`:

```
String sql = "select stream sum(e.eventValue), e.groupValue from
su.interference.test.entity.Event e group by e.groupValue";

StreamQueue q = (StreamQueue) s.execute(sql);
while (true) {
    Object o = q.poll();
    ...
}
```

A necessary and important note: since the insertion can be carried out in several threads, the order in which records are analyzed and grouped is based on the value of the identifier column (@Id), therefore, we strongly recommend using the @GeneratedValue annotation for the identifier, which ensures that the order of the increment identifier.

## Sliding windows

In this case, unlike the previous case, the output grouped record is generated for each newly inserted record, and the calculation of group function values is carried out for some group of records, the size of which is determined using the keyword WINDOW BY column INTERVAL = value. It should be noted that the syntax of this keyword differs from the generally accepted one - the only column with values and the interval by which the window size is determined are set directly in WINDOW BY:

```
String sql = "select stream count(e.eventId), sum(e.eventValue) from
su.interference.test.entity.Event e window by e.eventId interval =
100";

StreamQueue q = (StreamQueue) s.execute(sql);
while (true) {
    Object o = q.poll();
    ...
}
```

If the WINDOW BY keyword contains a column marked with @Id annotation, then the window size (in the rows) will be constant and equal to the value specified in INTERVAL.

## Distributed persistent model

To include a node in the cluster, you must specify the full list of cluster nodes (excluding this one) in the `cluster.nodes` configuration parameter. The minimum number of cluster nodes is 2, and the maximum is 64 (for more details, see cluster configuration rules below).

**Attention!** Cluster.nodes parameter must be filled completely before the first start of any of the nodes, and subsequently must not be changed. Then, you can start nodes in any order and at any time.

After such configuration, we may start all configured nodes as cluster. In this case, all nodes will be use specific messages (events) for provide inter-node data consistency and horizontal-scaling queries.

Interference cluster is a decentralized system. This means that the cluster does not use any coordination nodes; instead, each node follows to a set of some formal rules of behavior that guarantee the integrity and availability of data within a certain interaction framework.

Within the framework of these rules, all nodes of the Interference cluster are equivalent. This means that there is no separation in the system of master and slave nodes - changes to user tables can be made from any node, also all changes are replicated to all nodes, regardless of which node they were made on.

Running commit in a local user session automatically ensures that the changed data is visible on all nodes in the cluster.

## Distribute rules

- all cluster nodes are equivalent
- all changes on any of the nodes are mapped to other nodes
- if replication is not possible (the node is unavailable or the connection is broken), a persistent change queue is created for this node
- the owner of any data block is the node on which this block has been allocated
- the system uses the generation of unique identifiers for entities (@DistributedId) so that the identifier is uniquely unique within the cluster, and not just within the same node
- the system does not use any additional checks for uniqueness, requiring locks at the cluster level
- data inserts are performed strictly in local structures and then replicated
- changes (update / delete) can be performed only on the node-owner of the data block with this record. In future versions, it will be possible to change data from any node by obtaining a lock on the owner node.

## @DistributedId annotation

Using standard @Id, @GeneratedValue annotations implies the generation of unique values within a single node. If your distributed application is guaranteed that data insertion process will be performed on only one specific node, then this pair of annotations is enough. If the data can be inserted on different nodes, you must use the @DistributedId annotation with the above pair of annotations. This annotation guarantees the uniqueness of the generated identifier within the cluster and is highly recommended for use with @Id and @GeneratedValue.

## Distributed locks

When you try to change the object with a specific Id on any node at the cluster level, it is ensured that the object is locked on the owner node, i.e. on the node where this object was entered into the database. In this regard, it is recommended that mass changes be made on the data owner node, which will exclude additional remote locking requests. By default, the ability to change data from hosts that are not host hosts is disabled. You can enable it by setting the distribute.lock = true configuration parameter (Not available in version 2020.1). The object lock in the cluster is released by commit or rollback.

## Fault Tolerance

In the normal cluster operation mode, each of the nodes regularly sends messages to the remaining nodes of the cluster.

If the connection between the current node and a remote node is interrupted, all identifiers of the changed blocks that cannot be replicated to this node are saved until the functionality of this node is restored.

After restoration of the node's operability, the node checks the availability of all other nodes. The remaining nodes, in turn, roll forward to the offline node of the changes that occurred during its inoperability. After the roll of changes and subject to the availability of all other nodes, the node goes online.

## SQL horizontal-scaling queries

All SQL queries called on any of the cluster nodes will be automatically distributed among the cluster nodes for parallel processing, if such a decision is made by the node based on the analysis of the volume of tasks (the volume of the query tables is large enough, etc.)

If during the processing of a request a node is unavailable, the task distributed for this node will be automatically rescheduled to another available node.

## Configuration parameters

The current configuration is contained in the `config/interference.properties` file

The following describes the values of the configuration parameters and provides the optimal values for most applications of the parameters:

**local.node.id** - node identifier in the cluster - integer value from 1 to 64. All nodes in the cluster must have unique identifiers. The parameter must be specified when creating the instance and cannot be changed further.

**files.amount** - the number of threads that have the ability to simultaneously execute changes to the repository. Each thread operates with own, unique selected file. It is recommended to set a value equal to the number of processor cores. The default value is 4. Values from 1 to 64 can be used. The parameter must be specified when creating the instance and cannot be changed further.

**frame.size** is the size of the physical storage block. The default value is 8192. Values from 2048 to 65536 can be used. The parameter must be specified when creating the instance and cannot be changed further.

**frame.size.ix** - the size of the physical block for storing indexes. The default is 4096. Upgrading is not recommended. Values from 2048 to 65536 can be used. The parameter must be specified when creating the instance and cannot be changed further.

**codepage** - The codepage used to serialize string objects (String). The parameter must be specified when creating the instance and cannot be changed further.

**dateformat** - Date format used in SQL queries for string constants which used in WHERE clause condition, and, optionally in the management console (not use in 2020.1)

**db.path** - path to store data files

**journal.path** - path to store the journal file (not use in 2020.1)

**rmport** - initial numeric value which defining first server port for cluster transport interactions (see cluster configuration rules below).

**mmport** - http port for access to the control console via http protocol. The parameter must be specified when creating the instance and cannot be changed further.

**diskio.mode** - write mode to disk. 2 values are used - sync (write through mode) and async (write back mode). By default, sync is used and it is not recommended to change it.

**sync.period** - period of discard of changed blocks from the queue to disk in milliseconds. The default is 2000. For OLTP systems, it is recommended to set it to 100-1000, for storages with rare changes - at 5000-10000. Min value = 10, max value = 60000.

**sync.lock.enable** - lock data changes for the duration of a scheduled sync of blocks to disk. By default, set to true.

**cluster.nodes** - list of nodeIds, hosts and ports of cluster nodes, separated by commas. The list must contains string of the following format:

nodeId:host:port,nodeId:host:port, ... etc.

If the value is not set, the node will function in single mode (as local database)

**retrieve.threads.amount** - the number of threads for parallel processing of the SQL query.

**retrieve.queue.size** - size (amount of elements) of blocking queue, which use in SQL retrieve mechanism for prevent of heap overload. For best performance, use max possibly value depends of your heap size and amount of SQL queries running simultaneously. Usually, 10000-100000 is enough or optimal performance.

**auto.class.register** - a list of fully qualified names of entity classes, separated by commas, for which when the service starts, verification will be performed and, if necessary, automatic registration (both for services operating in standalone mode and at the application level)

**hbeat.period** - heartbeat period in milliseconds. The default is 1000.

## Cluster configuration rules

As we point above, interference cluster use specific messages for provide interaction between cluster nodes. Each any two nodes in cluster uses in each of two directions statically configured transport channel (client -> server), which provide message delivery from one node to other. Each transport channel is one-directional, so, we need two configured transport channels for full interaction between two nodes. Therefore, several event servers should run on each node, the number of which is equal to the amount of cluster nodes - 1 (or total amount of parts of cluster.nodes parameter, which amount should be the same on each node). Rmport configuration parameter contains start value of server ports on current node, all additional server ports values calculated incrementally. In configuration this may described by next example (three nodes cluster with ip addresses = 192.168.100.1, 192.168.100.2, 192.168.100.3):

```
local.node.id=1
rmport=8050
cluster.nodes=2:192.168.100.2:8060,3:192.168.100.3:8070

local.node.id=2
rmport=8060
cluster.nodes=1:192.168.100.1:8050,3:192.168.100.3:8071

local.node.id=3
rmport=8070
cluster.nodes=1:192.168.100.1:8051,2:192.168.100.2:8061
```

As we see, each port may use only one time.

## Persistent data types

In the current version, Interference supports the following persistent data types, which can be used as data types of fields of Entity classes (transient fields can be use any type):

```
int
long
float
double
java.lang.Integer
java.lang.Long
java.lang.Float
java.lang.Double
java.util.concurrent.AtomicInteger
java.util.concurrent.AtomicLong
java.lang.String
java.util.Date
java.util.ArrayList using the above types
java.util.HashMap using the above types
[] using the above types
```